

Alex Chen

alex.chen@devmail.io • +1 (415) 867-5309 • [linkedin.com/in/alexchen-dev](https://www.linkedin.com/in/alexchen-dev) • github.com/alexchen-dev • alexchen.dev

EDUCATION

University of California, Berkeley

Aug 2018 – May 2022

B.S. Computer Science — GPA: 3.82 / 4.0

Relevant Coursework: Database Systems, Operating Systems, Computer Networks, Distributed Systems, Data Structures & Algorithms

EXPERIENCE

Stripe

Jul 2022 – Present

Software Engineer II — Payments Infrastructure · San Francisco, CA

- Designed and shipped a high-throughput idempotency layer handling 120K+ transactions/sec using Redis and Postgres, reducing duplicate charges by 99.97%.
- Re-architected the payment routing service as a gRPC microservice, cutting average API latency from 340ms to 58ms (83% reduction).
- Authored internal SDK for 3 partner banks integrating Stripe Connect, abstracting OAuth 2.0 flows and webhook verification.
- Built real-time fraud signal aggregation pipeline using Kafka + Flink, feeding 40+ ML model features with <200ms end-to-end latency.

Lyft

May 2021 – Aug 2021

Data Engineering Intern · Remote

- Developed an automated ELT pipeline in Apache Airflow + dbt to ingest and model 8 TB/day of ride-event data into Snowflake.
- Reduced Spark job runtime by 62% via partition pruning and broadcast join optimizations across 15 production DAGs.
- Containerized all Airflow tasks using Docker and deployed via Kubernetes on GKE.

Accenture

Jun 2020 – Aug 2020

Technology Analyst (Frontend) · Remote

- Integrated GraphQL subscriptions for real-time dashboard updates, eliminating 12 REST polling endpoints.
- Implemented responsive layouts and WCAG 2.1 AA accessibility improvements, reducing reported a11y issues by 78%.

PROJECTS

DistribuCachE

github.com/alexchen-dev/distribucache

Tech: Go, gRPC, Protocol Buffers, Redis, Docker

A consistent-hashing distributed cache in Go supporting LRU eviction, replication factor 3, and automatic node rebalancing. Benchmarked at 800K ops/sec on 4-node cluster.

NeuralResume

github.com/alexchen-dev/neural-resume

Tech: Python, PyTorch, HuggingFace, FastAPI, PostgreSQL

Fine-tuned a BERT-based NLP classifier on 12K labeled job descriptions to extract skills and match candidates. Achieved F1=0.91 on held-out test set. Packaged as a FastAPI microservice.

K8s-Autopilot

github.com/alexchen-dev/k8s-autopilot

Tech: Go, Kubernetes, Prometheus, Helm, Slack API

Kubernetes operator (Go) that watches custom CRDs and automatically scales deployments based on Prometheus metrics, with Slack alerting integration.

SKILLS

Languages: Python, Go, TypeScript / JavaScript, Java, SQL, Bash

Frameworks: FastAPI / Flask / Django, Apache Spark / Flink, Kafka Streams, gRPC / Protocol Buffers

Tools: Docker / Kubernetes, Prometheus / Grafana, GitHub Actions / Jenkins, Git

Databases: PostgreSQL / MySQL, Redis, MongoDB, Elasticsearch

ACHIEVEMENTS & CERTIFICATIONS

- AWS Certified Solutions Architect – Associate

Amazon Web Services 2023

- HackMIT 2021 — Best Infrastructure Hack

MIT 2021

- Speaker: PyCon US 2023 — 'Building Fraud Signals at Scale with Kafka and Python'

Python Software Foundation 2023

- UC Berkeley Dean's List (6 semesters)

UC Berkeley 2018–2022